

Absolutely! Let's simplify those formulas into **plain-English-style math** so it's easier to understand. No fancy symbols, just **clear, step-by-step logic** using basic notation. We'll assume the image is grayscale with pixel values between **0 and 255** (like most 8-bit images).

 **Goal: Spread out pixel brightness values to improve image contrast.**

Simplified Step-by-Step with Easy Formulas

Step 1: Build the Histogram

Count how many times each pixel value (0–255) appears in the image.

$\text{hist}[i] = \text{number of pixels with value } i \text{ (for } i = 0 \text{ to } 255)$

You end up with an array `hist[256]`.

Step 2: Normalize the Histogram (PDF)

Divide each value in `hist[]` by total number of pixels in the image.

$\text{pdf}[i] = \text{hist}[i] / \text{total_pixels}$

This tells you the **probability** of each pixel value.

Step 3: Calculate CDF (Cumulative Distribution Function)

For each value from 0 to 255, add up all pdf values from 0 up to that index.

$\text{cdf}[0] = \text{pdf}[0]$

$\text{cdf}[1] = \text{cdf}[0] + \text{pdf}[1]$

$\text{cdf}[2] = \text{cdf}[1] + \text{pdf}[2]$

...

`cdf[i] = cdf[i - 1] + pdf[i]`

You get another array `cdf[256]`.

Step 4: Create Mapping for New Pixel Values

Multiply each value in `cdf[]` by 255 (max pixel value), and round it to the nearest integer.

`new_value[i] = round(cdf[i] * 255)`

Now you have a mapping: original value `i` should be changed to `new_value[i]`.

Step 5: Apply New Values to the Image

Go through every pixel in the image, and replace it:

`output_pixel = new_value[original_pixel]`

This gives you the equalized image!



Example with Tiny Image

Assume image: `[52, 55, 52, 60, 55]`

- Histogram: count how many times each number appears
 - PDF: divide counts by 5 (total pixels)
 - CDF: keep adding from left to right
 - Map each original value to a new value
 - Replace values in the image using the mapping
-



Want Even Simpler?

Just remember:

- Count how often each brightness happens → `hist`
 - Turn that into % probability → `pdf`
 - Add up the % → `cdf`
 - Multiply by 255 → get new values
 - Replace old values → done!
-

Let me know if you want this in code (Python, C++, or Java), or with a real image example.

Code Explanation (CUDA Implementation)

Let's go back to the CUDA code provided earlier. Here's how it works in detail, specifically for histogram equalization:

1. Loading the Image:

- The image is loaded and converted to grayscale using OpenCV. Grayscale images are easier to process because each pixel has a single intensity value (between 0 and 255).

2. CUDA Memory Allocation:

- Memory is allocated on the GPU for the image, histogram, CDF, and output image.

3. Compute Histogram (CUDA Kernel):

- Each pixel's intensity is checked, and the histogram is updated. The `atomicAdd` function ensures that multiple threads do not conflict when updating the same histogram bin.

4. Compute CDF (CUDA Kernel):

- The histogram is used to compute the cumulative sum, creating the CDF.

5. Equalization (CUDA Kernel):

- The CDF is then used to map each pixel in the image to a new intensity value, effectively equalizing the histogram.

6. Final Steps:

- The equalized image is copied back to the CPU and saved.

- A comparison of pixel intensities before and after equalization helps to understand the improvement in contrast.
-

Conclusion

Histogram equalization is a simple and effective technique to improve the contrast of an image, making features more visible. The basic steps involve computing a histogram, calculating a CDF, normalizing it, and then mapping the pixel values to the new range. While it's straightforward, it's a powerful tool for image enhancement, and implementing it using CUDA can significantly speed up the process, especially for large images.

Let me know if you need more clarification on any part!

Code Walkthrough: Step-by-Step

Step 1: Load and Resize Image

```
cv::Mat img = cv::imread("Sandesh.jpeg", cv::IMREAD_GRAYSCALE);  
cv::resize(img, img, cv::Size(1024, 1024));
```

- Loads a grayscale image.
 - Resizes it to $1024 \times 1024 = 1,048,576$ pixels for consistency and performance.
-

Step 2: Allocate GPU Memory

```
cudaMalloc(&d_in, total);  
cudaMalloc(&d_out, total);  
cudaMalloc(&d_hist, SIZE * sizeof(int));  
cudaMalloc(&d_cdf, SIZE * sizeof(int));
```

- Allocate memory on GPU for:
 - **d_in**: input image
 - **d_out**: output image
 - **d_hist**: histogram
 - **d_cdf**: cumulative distribution function

Step 3: Copy Image to GPU and Initialize Histogram

```
cudaMemcpy(d_in, img.data, total, cudaMemcpyHostToDevice);  
cudaMemset(d_hist, 0, SIZE * sizeof(int));
```

- Copy image to GPU.
- Set histogram to all zeros.

Step 4: Start Timer

```
clock_t start = clock();
```

- To calculate GPU execution time.

Step 5: CUDA Kernel for Histogram Calculation

```
computeHist<<<(total + 255) / 256, 256>>>(d_in, d_hist, total);
```

- Each thread checks a pixel and updates the `hist` using `atomicAdd` (because multiple threads can access same bin).

Step 6: CUDA Kernel for CDF Calculation

```
computeCDF<<<1, 1>>>(d_hist, d_cdf, total);
```

- Runs a single thread block to:
 - Compute cumulative histogram.
 - Normalize to range [0, 255].

Step 7: Apply Equalization Using CUDA

```
equalize<<<(total + 255) / 256, 256>>>(d_in, d_out, d_cdf, total);
```

- Each thread replaces a pixel's value based on the `cdf`.
-

Step 8: Stop Timer & Print Time

```
clock_t end = clock();  
double duration = double(end - start) / CLOCKS_PER_SEC * 1000;
```

Step 9: Copy Equalized Image Back to Host

```
cudaMemcpy(result.data, d_out, total, cudaMemcpyDeviceToHost);
```

Step 10: Analyze Intensities Before & After Equalization

```
map<int, int> oldIntensities, newIntensities;  
for (...) oldIntensities[img.data[i]]++;  
for (...) newIntensities[result.data[i]]++;
```

- Count how many pixels existed for each intensity (0–255) in:
 - Original image.
 - Equalized image.
 - Helps verify that the intensity range has been equalized (contrast improved).
-

Step 11: Save Result and Free Memory

```
cv::imwrite("result_GPU.jpg", result);  
cudaFree(...);
```

- Save the equalized image.
 - Free all GPU memory.
-

Summary

This CUDA-based program speeds up **histogram equalization**, a key image enhancement technique, by parallelizing:

- Histogram computation
- CDF computation
- Pixel transformation

Let me know if you want a **CPU vs GPU performance comparison**, **visual plot of histograms**, or how to implement the **CPU version** for comparison.