

Let's take a **2x2 matrix** to explain LU Decomposition step-by-step.

Example:

We have the matrix A :

$$A = \begin{bmatrix} 4 & 3 \\ 6 & 5 \end{bmatrix}$$

We want to decompose this matrix A into two matrices: L (Lower triangular matrix) and U (Upper triangular matrix).

Step 1: Initialize L and U

We know that in LU Decomposition:

- L is a lower triangular matrix with 1's on the diagonal.
- U is an upper triangular matrix.

At the start, we initialize:

$$L = \begin{bmatrix} 1 & 0 \\ l_{21} & 1 \end{bmatrix} \quad \text{and} \quad U = \begin{bmatrix} u_{11} & u_{12} \\ 0 & u_{22} \end{bmatrix}$$

We will now compute the values of l_{21} , u_{11} , u_{12} , and u_{22} .

Step 2: Start with the first row and column

Calculate u_{11} , u_{12} :

- The first element of U , u_{11} , is the same as the first element of matrix A :

$$u_{11} = 4$$

- The next element in the first row of U , u_{12} , is the second element of the first row of matrix A :

$$u_{12} = 3$$

At this point, we have:

$$L = \begin{bmatrix} 1 & 0 \\ l_{21} & 1 \end{bmatrix} \quad \text{and} \quad U = \begin{bmatrix} 4 & 3 \\ 0 & u_{22} \end{bmatrix}$$

Step 3: Move to the second row

Calculate l_{21} and u_{22} :

- We now compute l_{21} , which will be used to fill the second row of L .

From the formula:

$$l_{21} \cdot u_{11} = a_{21} \quad l_{21} \cdot u_{11} = a_{21}$$

Substitute the known values:

$$l_{21} \cdot 4 = 6 \Rightarrow l_{21} = \frac{6}{4} = 1.5 \quad l_{21} \cdot 4 = 6 \quad \Rightarrow \quad l_{21} = \frac{6}{4} = 1.5$$

Now, for the second element of U , we compute u_{22} :

$$u_{22} = a_{22} - l_{21} \cdot u_{12} \quad u_{22} = a_{22} - l_{21} \cdot u_{12}$$

Substitute the known values:

$$u_{22} = 5 - (1.5 \cdot 3) = 5 - 4.5 = 0.5 \quad u_{22} = 5 - (1.5 \cdot 3) = 5 - 4.5 = 0.5$$

At this point, the matrices are:

$$L = \begin{bmatrix} 1 & 0 \\ 1.5 & 1 \end{bmatrix} \quad \text{and} \quad U = \begin{bmatrix} 4 & 3 \\ 0 & 0.5 \end{bmatrix}$$

Final Result:

Thus, the LU decomposition of the matrix A is:

$$A = \begin{bmatrix} 4 & 3 \\ 6 & 5 \end{bmatrix} = L \cdot U = \begin{bmatrix} 1 & 0 \\ 1.5 & 1 \end{bmatrix} \cdot \begin{bmatrix} 4 & 3 \\ 0 & 0.5 \end{bmatrix}$$

Verification:

To verify, we can multiply the L and U matrices:

$$L \cdot U = \begin{bmatrix} 1 & 0 \\ 1.5 & 1 \end{bmatrix} \cdot \begin{bmatrix} 4 & 3 \\ 0 & 0.5 \end{bmatrix} = \begin{bmatrix} 4 & 3 \\ 6 & 5 \end{bmatrix}$$

The result is exactly the original matrix A , confirming that the LU decomposition is correct.

Summary:

- **LU Decomposition** decomposes a matrix A into a lower triangular matrix L and an upper triangular matrix U .
- In the decomposition, L contains the multipliers (coefficients) used to eliminate elements in A , and U contains the resulting upper triangular matrix.

Example of LU Decomposition using Gaussian Elimination for a 2x2 Matrix

Let's take the matrix A as an example:

$$A = \begin{bmatrix} 2 & 3 \\ 4 & 7 \end{bmatrix}$$

We will decompose this matrix into L (lower triangular matrix) and U (upper triangular matrix).

Step 1: Initialize Matrices

We start by initializing the matrices:

- L is initialized as the identity matrix.
- U is initialized as a copy of A .

$$L = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad U = \begin{bmatrix} 2 & 3 \\ 4 & 7 \end{bmatrix}$$

Step 2: Eliminate Elements Below the Pivot

We perform Gaussian elimination to make the elements below the pivot (the element $U[1][1] = 2$) zero.

For row 2 (element $U[2][1] = 4$):

We calculate the multiplier m :

$$m = \frac{U[2][1]}{U[1][1]} = \frac{4}{2} = 2$$

Now, we update row 2 of matrix U :

$$U[2] = U[2] - m \times U[1] = [4, 7] - 2 \times [2, 3] = [4, 7] - [4, 6] = [0, 1]$$

We also update the L matrix with the multiplier:

$$L[2][1] = m = 2$$

So after this step, the matrices L and U are:

$$L = \begin{bmatrix} 1 & 0 \\ 2 & 1 \end{bmatrix} \quad U = \begin{bmatrix} 2 & 3 \\ 0 & 1 \end{bmatrix}$$

Step 3: Final Result

At this point, the matrix A has been decomposed into L and U such that:

$$A = L \cdot U$$

So the final result is:

$$L = \begin{bmatrix} 1 & 0 \\ 2 & 1 \end{bmatrix} \quad U = \begin{bmatrix} 2 & 3 \\ 0 & 1 \end{bmatrix}$$

Verification

Let's multiply LL and UU to verify that it equals AA:

$$L \cdot U = \begin{bmatrix} 1 & 0 \\ 2 & 3 \end{bmatrix} \cdot \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 2 & 1 \\ 4 & 7 \end{bmatrix} = A$$

Thus, the LU decomposition is correct.

Conclusion

The LU decomposition of matrix AA is:

$$L = \begin{bmatrix} 1 & 0 \\ 2 & 3 \end{bmatrix} \quad U = \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix}$$

This process shows how we can decompose a 2x2 matrix into its lower and upper triangular matrices using **Gaussian Elimination**.

Comparison of Gaussian Elimination and LU Decomposition

Both **Gaussian Elimination** and **LU Decomposition** are techniques used to solve systems of linear equations, find the determinant, and invert matrices. However, they have distinct purposes, and their use cases, implementation, and computational complexities differ.

Here's a comparison of **Gaussian Elimination** and **LU Decomposition**:

Aspect	Gaussian Elimination	LU Decomposition
Purpose	Solving a system of linear equations $Ax=b$ directly.	Factorizing a matrix A into a product of lower and upper triangular matrices L and U .
Matrix Form	Used for solving a system of equations for a matrix A and vector b .	Decomposes matrix A into L (lower triangular) and U (upper triangular).
Methodology	Successive row operations to eliminate variables and solve for unknowns.	Factorizes A once into L and U , then solves using forward and backward substitution.
Steps	<ul style="list-style-type: none">- Forward elimination to create upper triangular matrix U.- Back substitution to solve for unknowns.	<ul style="list-style-type: none">- Perform Gaussian elimination to get L and U.- Use L and U to solve the system using forward and backward substitution.
Efficiency	Efficient for a single system of equations.	Efficient for multiple systems of equations (solve $Ax=b$ for many different b vectors after LU decomposition).
Numerical Stability	Susceptible to round-off errors, especially for ill-conditioned matrices.	More stable, especially if using pivoting (partial or complete).
Pivoting	Requires pivoting for stability, especially when the pivot element is small.	Pivoting is often used, but it's done explicitly in the LU factorization step for numerical stability.

Usage	Solves a system of equations directly.	Decomposes a matrix once, making it easier to solve many systems with the same matrix A .
Determinant Calculation	Determinant can be computed during the elimination process by multiplying the pivots.	Determinant is the product of the diagonal elements of UU (i.e., the pivots of the LU decomposition).
Inversion	Matrix inversion requires solving for each column vector one at a time.	Matrix inversion is easier after LU decomposition: use forward and backward substitution for each column.

Computational Complexity Comparison

1. Gaussian Elimination:

- **Time Complexity:**

The computational complexity of Gaussian elimination is $O(n^3)$, where n is the size of the matrix (i.e., the number of rows or columns in an $n \times n$ matrix).

Why?

- The algorithm involves performing **forward elimination** to create an upper triangular matrix, which takes approximately $O(n^3)$ operations. This is done by performing row operations for each of the n rows and involving n elements per row.
- After that, the **back substitution** step is done in $O(n^2)$, but this is dominated by the forward elimination step.

- **Overall complexity:**

$O(n^3)$

2. LU Decomposition:

- **Time Complexity:**

The LU decomposition process also has a time complexity of $O(n^3)$.

Why?

- The LU decomposition is essentially a variant of Gaussian elimination. It performs similar row operations to decompose the matrix into a lower triangular matrix L and an upper triangular matrix U .

- This process takes $O(n^3)$ because we still perform row operations on each element of the matrix, just as in Gaussian elimination.
- **After LU decomposition**, solving the system of equations using forward and backward substitution takes $O(n^2)$ operations (because you solve for each of the n variables using the n -by- n triangular matrices).

Overall complexity for LU decomposition and solving once:

$O(n^3)$ (for decomposition) + $O(n^2)$ (for solving) = $O(n^3)$ \text{ (for decomposition)} + $O(n^2)$ \text{ (for solving)} = $O(n^3)$

Memory Requirements:

1. **Gaussian Elimination:**
 - The algorithm requires $O(n^2)$ space to store the matrix A and vector b , because you store the entire augmented matrix in memory.
2. **LU Decomposition:**
 - The algorithm requires $O(n^2)$ space for storing both the original matrix A and the resulting L and U matrices, since both L and U are $n \times n$ \times n matrices.
3. In some cases, if you perform the decomposition **in-place**, you might only need $O(n^2)$ space for storing the matrix A (overwritten with L and U).

Summary Table:

Aspect	Gaussian Elimination	LU Decomposition
Time Complexity	$O(n^3)$	$O(n^3)$
Space Complexity	$O(n^2)$	$O(n^2)$
Best Use Case	Solving one system of equations.	Solving multiple systems with the same matrix.
Numerical Stability	Less stable without pivoting.	More stable with pivoting.

Efficiency

Direct solution, one system at a time.

Efficient for multiple systems.

Conclusion:

- **Gaussian elimination** is best suited for solving a **single system of linear equations** where you need the result right away.
- **LU decomposition** is more efficient when solving **multiple systems of equations** that share the same coefficient matrix, as you only need to perform the decomposition once.

LU Decomposition vs LUP Decomposition

Both **LU decomposition** and **LUP decomposition** are methods used to factorize a square matrix A into triangular matrices L (lower triangular matrix) and U (upper triangular matrix). The difference between **LU** and **LUP** comes in handling matrix **pivoting**, which is critical for numerical stability, especially when dealing with ill-conditioned or singular matrices.

Here's a detailed comparison of **LU decomposition** and **LUP decomposition**:

LU Decomposition (without Pivoting)

Description:

- LU decomposition factors a matrix A into the product of a lower triangular matrix L and an upper triangular matrix U , such that:
 $A = LU$
- It assumes that the matrix is **non-singular** (invertible) and that no row exchanges are needed.

Steps:

1. The matrix A is decomposed into L and U directly through **Gaussian elimination**.
2. The diagonal elements of L are 1, and U is upper triangular.
3. Solving $Ax = b$ is done by first solving $Ly = b$ (forward substitution) and then $Ux = y$ (backward substitution).

Limitations:

- **Numerical Stability:** LU decomposition can be numerically unstable when the matrix has very small pivot elements (i.e., it is ill-conditioned). This could lead to **large errors** in the computation.
 - **Pivoting:** LU decomposition does **not** use pivoting, which means if a small pivot element is encountered, the results could be inaccurate or invalid.
-

LUP Decomposition (LU Decomposition with Partial Pivoting)

Description:

- **LUP decomposition** is an enhanced version of LU decomposition that **incorporates pivoting** to improve numerical stability.
- It decomposes the matrix A into three matrices:
 $A = LUP$
where:

- LL is a lower triangular matrix with unit diagonal elements.
- UU is an upper triangular matrix.
- PP is a **permutation matrix** that keeps track of row exchanges.

Steps:

1. **Partial Pivoting:** Before each elimination step, the row with the largest absolute value in the pivot column is selected and swapped with the current row (this helps avoid division by small numbers).
2. The matrix AA is decomposed into LL, UU, and PP via **Gaussian elimination with partial pivoting**.
3. To solve $Ax=b$, the system is modified to $LUx=Pb$. The solution is obtained by solving:
 - $Ly=Pb$ (forward substitution)
 - $Ux=y$ (backward substitution)

Advantages:

- **Numerical Stability:** LUP decomposition is **more stable** than LU decomposition because the pivoting process ensures that the largest available elements are used as pivots during the elimination process. This reduces the risk of numerical instability in the computation.
- **Works with Singular Matrices:** If a matrix is **singular** or near-singular (i.e., non-invertible or ill-conditioned), LUP decomposition can still provide meaningful results by handling row exchanges properly. It can also detect whether the matrix is singular by checking if any of the diagonal elements of UU become zero.

Limitations:

- The addition of the **permutation matrix** introduces an extra step (storing the row swaps), but this does not affect the overall complexity.

Comparison Table:

Aspect	LU Decomposition	LUP Decomposition
Pivoting	No pivoting is performed.	Pivoting (partial) is used to improve stability.

Matrix Factorization	$A=LU$	$A=LUP$
Numerical Stability	Can be unstable for ill-conditioned matrices.	More stable due to row swaps (pivoting).
Complexity	$O(n^3)$ for the decomposition.	$O(n^3)$ for the decomposition.
Use Case	Suitable for matrices that are well-conditioned and don't require row swapping.	Used for matrices that may have small pivots or are ill-conditioned.
Matrix Size	L and U are both $n \times n$.	L , U , and P are all $n \times n$.
Inversion	Requires solving for each column of the inverse.	Similar, but with an additional step for row permutation.
Singularity Detection	Does not handle singular matrices well.	Can detect singularity if any pivot is zero.

Computational Complexity:

Both **LU** and **LUP decomposition** have the same computational complexity:

- **Time Complexity:**
 - Both algorithms involve $O(n^3)$ operations for decomposition, where n is the size of the matrix (i.e., $n \times n$).
- **Space Complexity:**
 - Both LU and LUP decomposition require $O(n^2)$ space to store the matrices L , U , and P (in the case of LUP).

Example:

Given Matrix A (for LUP Decomposition):

$$A = \begin{bmatrix} 2 & 3 \\ 5 & 4 \end{bmatrix}$$

LU Decomposition:

Without pivoting, we would perform Gaussian elimination on the matrix, but if the pivot element is very small (in this case, $A[1][1] = 2$), it can cause numerical issues.

LUP Decomposition (with pivoting):

The process will swap rows to make sure the pivot element is the largest available in the column. After pivoting, the matrix might look like this:

$$P = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, L = \begin{bmatrix} 1 & 0 \\ 0.4 & 1 \end{bmatrix}, U = \begin{bmatrix} 5 & 4 \\ 0 & 0.6 \end{bmatrix}$$

Conclusion:

- **LU Decomposition** is suitable when the matrix is well-conditioned and no pivoting is needed.
- **LUP Decomposition** is preferred for matrices where pivoting is necessary for numerical stability, especially for ill-conditioned or singular matrices. It guarantees better stability and helps handle row swaps efficiently.

Applications of LU Decomposition

LU Decomposition is widely used in numerical linear algebra, engineering, computer science, and various scientific fields. Here are some common applications of **LU decomposition**:

1. Solving Linear Systems

LU Decomposition is commonly used to solve systems of linear equations of the form:

$$Ax = b$$

where A is a square matrix, x is the solution vector, and b is the right-hand side vector. LU decomposition simplifies this process by breaking down A into two matrices L (lower triangular) and U (upper triangular), so that $A = LU$. This enables us to solve the system in two steps:

- Solve $Ly = b$ ($y = b$) (forward substitution)
- Solve $Ux = y$ ($x = y$) (backward substitution)

2. Matrix Inversion

LU decomposition can be used to compute the inverse of a matrix. By performing LU decomposition on matrix A , you can solve for each column of the inverse matrix in the form $A^{-1}x = e_i$, where e_i is the standard basis vector (a column vector with a 1 in the i -th position and 0s elsewhere). This is useful in computational applications such as simulation, optimization, and control systems.

3. Determinant Calculation

The determinant of a matrix A can be easily calculated from its LU decomposition. If $A = LU$, then the determinant of A is the product of the diagonal elements of U (since $\det(L) = 1$ for lower triangular matrix L):

$$\det(A) = \prod_{i=1}^n u_{ii}$$

where u_{ii} are the diagonal elements of U .

4. Eigenvalue Problems

LU decomposition is sometimes used as a preprocessing step in methods like **LU-based eigenvalue solvers** or **iterative methods** for finding eigenvalues and eigenvectors of a matrix.

5. Numerical Methods (Finite Element Methods)

In computational fluid dynamics, structural engineering, and other fields involving **finite element methods (FEM)**, LU decomposition is used to solve large systems of linear equations that arise from discretizing differential equations. FEM discretizes the governing equations into a large sparse matrix, and LU decomposition helps in efficiently solving these systems.

6. Optimization Problems

In optimization, LU decomposition is frequently used to solve large systems of equations that arise in **linear programming**, **least squares problems**, and other optimization techniques, where solving $Ax = b$ is required.

Example of Solving Linear Equation Using LU Decomposition

Let's take a **2x2 matrix** and solve a system of linear equations using **LU Decomposition**.

Problem:

Solve the system of linear equations:

$$3x+4y=7 \quad 2x+3y=5$$

This can be written in matrix form as:

$$A = \begin{bmatrix} 3 & 4 \\ 2 & 3 \end{bmatrix}, \quad b = \begin{bmatrix} 7 \\ 5 \end{bmatrix}$$

We want to solve for $\begin{bmatrix} x \\ y \end{bmatrix}$ (i.e., the vector $\begin{bmatrix} x \\ y \end{bmatrix}$).

Step 1: LU Decomposition of A

Perform LU decomposition on A:

$$A = \begin{bmatrix} 3 & 4 \\ 2 & 3 \end{bmatrix}$$

We decompose A into lower triangular matrix L and upper triangular matrix U:

$$A = LU \Rightarrow L = \begin{bmatrix} 1 & 0 \\ \frac{2}{3} & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 3 & 4 \\ 0 & \frac{4}{3} \end{bmatrix}$$

Step 2: Solve $Ly = b$ (Forward Substitution)

Now, solve the intermediate system $Ly = b$:

$$\begin{bmatrix} 1 & 0 \\ \frac{2}{3} & 1 \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \end{bmatrix} = \begin{bmatrix} 7 \\ 5 \end{bmatrix}$$

This results in the following system of equations:

$$y_1 = 7 \quad \frac{2}{3}y_1 + y_2 = 5 \Rightarrow 2 \times 7 + y_2 = 5 \Rightarrow 4.67 + y_2 = 5 \Rightarrow y_2 = 0.33$$

So, the vector y is:

$$y = \begin{bmatrix} 7 \\ 0.33 \end{bmatrix}$$

Step 3: Solve $Ux = y$ (Backward Substitution)

Now, solve the upper triangular system $Ux = y$:

$$\begin{bmatrix} 3 & 4 \\ 0 & \frac{4}{3} \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 7 \\ 0.33 \end{bmatrix}$$

This gives the following system of equations:

$$3x + 4y = 7 \quad 3x + 4y = 7 \quad 4y = 0.33 \frac{4}{3} y = 0.33$$

Solve for y from the second equation:

$$y = 0.33 \times 3 = 0.99 \quad y = \frac{0.33 \times 3}{4} = 0.25$$

Substitute $y = 0.25$ into the first equation:

$$3x + 4(0.25) = 7 \Rightarrow 3x + 1 = 7 \Rightarrow 3x = 6 \Rightarrow x = 2 \quad 3x + 4(0.25) = 7 \quad \Rightarrow \quad 3x + 1 = 7 \quad \Rightarrow \quad 3x = 6 \quad \Rightarrow \quad x = 2$$

So, the solution to the system of equations is:

$$x = 2, y = 0.25 \quad x = 2, \quad y = 0.25$$

Summary of Steps:

1. **LU Decompose** the matrix A into L and U .
2. **Solve** $Ly = b$ $y = b$ (forward substitution) to find the intermediate vector y .
3. **Solve** $Ux = y$ $x = y$ (backward substitution) to find the final solution vector x .

Computational Complexity:

- **LU Decomposition:**
 - Time complexity: $O(n^3)$, where n is the size of the matrix.
 - Space complexity: $O(n^2)$, for storing the matrices L and U .
- **Solving the Linear System:**
 - After LU decomposition, solving $Ly = b$ and $Ux = y$ takes $O(n^2)$ time (since forward and backward substitution are linear in complexity).

Thus, LU decomposition is efficient for solving multiple linear systems with the same matrix A , as the decomposition step is performed only once.

This C++ CUDA code implements LU Decomposition using GPU. LU decomposition is a method for decomposing a matrix into a product of a lower triangular matrix LL and an upper triangular matrix UU, such that:

$$A=LU = LU$$

Code Overview

Global Variables and Constants

```
const int N = 1024;
```

```
const int BLOCK_SIZE = 256;
```

- **N**: Size of the square matrix AA (i.e., $N \times N$ times N).
- **BLOCK_SIZE**: The number of threads per block for CUDA kernel execution. The block size is set to 256 threads.

CUDA Kernels

The kernels `compute_u` and `compute_l` are defined to calculate the upper matrix UU and lower matrix LL respectively.

1. `compute_u` Kernel

```
__global__ void compute_u(float* A, float* L, float* U, int i) {  
  
    int col = blockIdx.x * blockDim.x + threadIdx.x + i;  
  
    if (col < N) {  
  
        float sum = 0;  
  
        for (int j = 0; j < i; j++) {  
  
            sum += L[i * N + j] * U[j * N + col];  
  
        }  
  
        U[i * N + col] = A[i * N + col] - sum;  
  
    }  
}
```

- **Purpose**: This kernel computes the elements of the upper triangular matrix UU (for row ii).
- **Details**:

- Each thread processes a column element of the matrix.
- The sum is computed using elements from the lower matrix LL and the upper matrix UU, and the corresponding value in UU is updated.
- The formula used here is:

$$U[i,j] = A[i,j] - \sum_{k=0}^{i-1} L[i,k] \cdot U[k,j]$$

$$U[i, j] = A[i, j] - \sum_{k=0}^{i-1} L[i, k] \cdot U[k, j]$$

2. compute_l Kernel

```
__global__ void compute_l(float* A, float* L, float* U, int i) {
    int row = blockIdx.x * blockDim.x + threadIdx.x + i + 1;

    if (row < N) {
        float sum = 0;

        for (int j = 0; j < i; j++) {
            sum += L[row * N + j] * U[j * N + i];
        }

        L[row * N + i] = (A[row * N + i] - sum) / U[i * N + i];
    }
}
```

- **Purpose:** This kernel computes the elements of the lower triangular matrix LL (for column ii).
- **Details:**
 - Each thread processes a row element of the matrix.
 - The sum is computed using elements from the lower matrix LL and the upper matrix UU, and the corresponding value in LL is updated.
 - The formula used here is:

$$L[j,i] = \frac{A[j,i] - \sum_{k=0}^{i-1} L[j,k] \cdot U[k,i]}{U[i,i]}$$

$$L[j, i] = \frac{A[j, i] - \sum_{k=0}^{i-1} L[j, k] \cdot U[k, i]}{U[i, i]}$$

Matrix Initialization

A = new float[N * N];

L = new float[N * N];

U = new float[N * N];

- **A**: A dynamically allocated matrix representing the original square matrix AA.
- **L** and **U**: These are matrices to store the lower and upper triangular matrices, respectively.

Random Matrix Generation

```
srand(time(NULL));

for (int i = 0; i < N * N; i++) {

    A[i] = (float)(rand() % 1000) / 10.0f;

    L[i] = 0.0f;

    U[i] = 0.0f;

}
```

- **Random Values for Matrix AA**: The elements of matrix AA are randomly generated between 0 and 100. These values are stored in a 1D array.
- **Initialize Matrices LL and UU**: LL and UU are initialized to zeros.

Set LL as Identity Matrix

```
for (int i = 0; i < N; i++) {

    L[i * N + i] = 1.0f;

}
```

- **Identity Matrix**: The diagonal elements of LL are set to 1. This is a property of the lower triangular matrix in LU decomposition.

Memory Allocation on the Device

```
cudaMalloc(&d_A, N * N * sizeof(float));

cudaMalloc(&d_L, N * N * sizeof(float));

cudaMalloc(&d_U, N * N * sizeof(float));
```

- **Memory Allocation on GPU**: The device memory is allocated for the matrices AA, LL, and UU using `cudaMalloc`.

Copying Data to GPU

```
cudaMemcpy(d_A, A, N * N * sizeof(float), cudaMemcpyHostToDevice);  
  
cudaMemcpy(d_L, L, N * N * sizeof(float), cudaMemcpyHostToDevice);  
  
cudaMemcpy(d_U, U, N * N * sizeof(float), cudaMemcpyHostToDevice);
```

- **Copy Data from Host to Device:** The randomly generated matrix AA, as well as the initialized matrices LL and UU, are copied from the host (CPU) to the device (GPU).

Timer Setup

```
cudaEvent_t start, stop;  
  
cudaEventCreate(&start);  
  
cudaEventCreate(&stop);  
  
cudaEventRecord(start);
```

- **Timer:** CUDA events are used to measure the execution time of the GPU computations.

Launch Kernels for LU Decomposition

```
for (int i = 0; i < N; i++) {  
  
    int blocks = (N - i + BLOCK_SIZE - 1) / BLOCK_SIZE;  
  
    compute_u<<<blocks, BLOCK_SIZE>>>(d_A, d_L, d_U, i);  
  
    if (i < N - 1)  
  
        compute_l<<<blocks, BLOCK_SIZE>>>(d_A, d_L, d_U, i);  
  
}
```

- **For Loop Over Rows:** This loop iterates over the rows and columns of the matrix and launches CUDA kernels to compute the UU and LL matrices.
 - **compute_u:** For each row, the kernel is called to compute the corresponding elements in the upper triangular matrix UU.
 - **compute_l:** For rows below the current row, the kernel is called to compute the corresponding elements in the lower triangular matrix LL.

Wait for Kernels to Finish

```
cudaEventRecord(stop);  
  
cudaEventSynchronize(stop);
```

- **Stop Timer:** The stop event is recorded and synchronized to ensure all GPU tasks are completed before measuring the time.

Copy Results from Device to Host

```
cudaMemcpy(L, d_L, N * N * sizeof(float), cudaMemcpyDeviceToHost);  
  
cudaMemcpy(U, d_U, N * N * sizeof(float), cudaMemcpyDeviceToHost);
```

- **Copy Data Back to Host:** After the LU decomposition is done on the GPU, the results (matrices LL and UU) are copied back to the host.

Print Execution Time

```
float milliseconds;  
  
cudaEventElapsedTime(&milliseconds, start, stop);  
  
cout << "\nCUDA LU decomposition time: " << milliseconds / 1000.0f << " seconds\n";
```

- **Execution Time:** The time taken for the LU decomposition is calculated using the difference between the start and stop events and printed in seconds.

Print Matrices

```
print_matrix("Original Matrix A", A, print_size);  
  
print_matrix("Lower Matrix L", L, print_size);  
  
print_matrix("Upper Matrix U", U, print_size);
```

- **Matrix Output:** The matrices AA, LL, and UU are printed (only the top 8x8 block of the matrix to avoid printing large matrices).

Clean Up

```
delete[] A;
```

```
delete[] L;
```

```
delete[] U;
```

```
cudaFree(d_A);
```

```
cudaFree(d_L);
```

```
cudaFree(d_U);
```

- **Memory Deallocation:** Both host and device memory are freed.
-

Explanation of Key Concepts

- **LU Decomposition:** It is a method of factorizing a matrix into two triangular matrices, LL (lower triangular) and UU (upper triangular), such that $A=LU$.
- **CUDA Kernels:** `compute_u` and `compute_l` are parallel functions running on the GPU to compute the elements of UU and LL.
- **Block and Grid Structure:** CUDA uses a grid of blocks with threads to parallelize the computation, speeding up the LU decomposition for large matrices.
- **Memory Management:** Efficient memory management using `cudaMalloc` and `cudaMemcpy` ensures that data is transferred between the host and the GPU efficiently.

Summary

This code performs LU decomposition on a randomly generated matrix using GPU, leveraging CUDA for parallel computing. The process involves breaking down the matrix into two triangular matrices LL and UU, solving for them efficiently using CUDA kernels, and measuring the performance.

This program performs **LU decomposition** (Lower-Upper decomposition) of a matrix AA, where AA is decomposed into two matrices: LL (Lower triangular matrix) and UU (Upper triangular matrix). The program uses **OpenACC** for parallel processing to speed up the computation.

Let's go through the program step by step:

1. Includes and Constants

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <time.h>
```

```
#define N 1024
```

- **stdio.h**: For input/output operations (like `printf`).
- **stdlib.h**: For memory allocation (`malloc`).
- **time.h**: For measuring the execution time of the LU decomposition process.
- **N**: A constant that defines the size of the matrix AA, LL, and UU as 1024×1024 .

2. Matrix Printing Function

```
void print_matrix(const char *name, float *mat, int size)
```

```
{  
    printf("\n%s (Top %dx%d block):\n", name, size, size);  
    for (int i = 0; i < size; i++)  
    {  
        for (int j = 0; j < size; j++)  
        {  
            printf("%8.2f ", mat[i * N + j]);  
        }  
        printf("\n");  
    }  
}
```

This function prints a **block of the matrix**. The `print_matrix` function prints the top left $\text{size} \times \text{size}$ block of the matrix to give an overview of the matrix content.

3. Memory Allocation and Initialization

```
float *A, *L, *U;  
  
A = (float *)malloc(N * N * sizeof(float));  
  
L = (float *)malloc(N * N * sizeof(float));  
  
U = (float *)malloc(N * N * sizeof(float));
```

```
srand(time(NULL));  
  
for (int i = 0; i < N * N; i++)  
{  
    A[i] = (float)(rand() % 1000) / 10.0f;  
    L[i] = 0.0f;  
    U[i] = 0.0f;  
}
```

```
for (int i = 0; i < N; i++)  
{  
    L[i * N + i] = 1.0f;  
}
```

- **Memory Allocation:** Allocates memory for three matrices: `A`, `L`, and `U`. Each of them is $N \times N$ in size.
 - `A`: Original matrix that will be decomposed.
 - `L`: Lower triangular matrix, initialized to zeros, except for the diagonal, which is set to 1 (since `LL` is a lower triangular matrix with ones on the diagonal).
 - `U`: Upper triangular matrix, initialized to zeros.

- **Random Initialization:** **A** is initialized with random values, and **L** and **U** are initialized with zero values.

4. OpenACC Parallelization Setup

The `#pragma acc` directives are used to instruct the compiler to parallelize parts of the code using **OpenACC**, which accelerates the computation by running loops in parallel on the GPU or multi-core CPU.

```
#pragma acc data copyin(A[0 : N * N]) copy(L[0 : N * N], U[0 : N * N])
```

```
{  
    for (int i = 0; i < N; i++)  
    {  
#pragma acc parallel loop  
        for (int j = i; j < N; j++)  
        {  
            float sum = 0.0f;  
            for (int k = 0; k < i; k++)  
            {  
                sum += L[i * N + k] * U[k * N + j];  
            }  
            U[i * N + j] = A[i * N + j] - sum;  
        }  
    }  
}
```

```
#pragma acc parallel loop  
    for (int j = i + 1; j < N; j++)  
    {  
        float sum = 0.0f;  
        for (int k = 0; k < i; k++)  
        {  
            sum += L[j * N + k] * U[k * N + i];  
        }  
    }
```



```

L[j * N + i] = (A[j * N + i] - sum) / U[i * N + i];
}
}
}

```

- **#pragma acc data:** This directive marks the section of the code that will benefit from parallelization and specifies which data arrays should be copied to and from the device (GPU).
 - **copyin:** Input data (from host to device).
 - **copy:** Input-output data (from host to device and from device to host).
- **Parallel Loops:**
 - The first parallel loop computes the upper triangular matrix UU.
 - The second parallel loop computes the lower triangular matrix LL.
- **The loops are structured as follows:**
 - The outer loop (**for (int i = 0; i < N; i++)**) iterates over the rows of AA to compute the LU decomposition.
 - The first nested loop computes the elements of UU.
 - The second nested loop computes the elements of LL.

Each of these loops is parallelized using the **#pragma acc parallel loop** directive.

5. Timing the LU Decomposition

```
clock_t start = clock();
```

```
// LU decomposition code here
```

```
clock_t end = clock();
```

```
double time_spent = (double)(end - start) / CLOCKS_PER_SEC;
```

```
printf("\nOpenACC LU decomposition time: %.6f seconds\n", time_spent);
```

The `clock()` function is used to measure the time taken to perform the LU decomposition. The difference between the start and end times is calculated and printed as the total time taken for LU decomposition.

6. Printing Matrices and Cleanup

```
int print_size = 8;

print_matrix("Original Matrix A", A, print_size);

print_matrix("Lower Matrix L", L, print_size);

print_matrix("Upper Matrix U", U, print_size);


free(A);

free(L);

free(U);
```

- After the decomposition, the program prints the top-left 8×8 blocks of matrices **A**, **L**, and **U**.
- The `free()` function is used to release the dynamically allocated memory for **A**, **L**, and **U**.

7. Compilation Command

```
pgc++ lup_openacc.cpp -O3 -Minfo=accel -o acc_lu && ./acc_lu
```

- **pgc++**: The PGI C++ compiler used for compiling the code with OpenACC support.
- **-O3**: Optimizes the code for maximum speed.
- **-Minfo=accel**: Provides information about the OpenACC directives and the parallelization performed.
- **-o acc_lu**: Specifies the output executable name.
- The executable is run with `./acc_lu`.

Key Points

- This program uses **OpenACC** for parallelism, speeding up the LU decomposition of large matrices by running matrix computations in parallel.
- The program performs LU decomposition by decomposing matrix AA into lower (LL) and upper (UU) triangular matrices.
- The performance is measured using the `clock()` function, and the parallel loops are executed in parallel using the `#pragma acc parallel loop` directives.

This program is a good example of using OpenACC to accelerate numerical algorithms (like LU decomposition) on GPUs or multi-core CPUs.

