

Sure! Let's go step by step — first, I'll explain **SYCL**, and then we'll break down your **Monte Carlo Pi estimation** program using SYCL.



What is SYCL?

SYCL (pronounced 'sickle') is a high-level programming model from the **Khronos Group** for **heterogeneous computing** (using both CPUs and GPUs).

It builds on **C++** and is designed to work well with **OpenCL**, enabling developers to write parallel code for **GPUs, CPUs, FPGAs**, etc., using **standard C++**.

Key concepts in SYCL:

- **queue**: Where you submit commands (kernels) to a device (GPU/CPU).
 - **buffer**: A memory object that automatically manages data movement between host and device.
 - **accessor**: Lets a kernel access memory inside a buffer.
 - **parallel_for**: Used to run code in parallel on many elements (threads).
 - SYCL is **single-source**: Both host and device code are written in the same C++ file.
-



What does this code do?

This is a **Monte Carlo simulation** to estimate the value of π (**pi**) using SYCL for parallel computation.



Step-by-step Breakdown:



1. Includes and Setup

```
#include <sycl/sycl.hpp>
#include <iostream>
#include <cmath>
#include <chrono>
```

- Brings in SYCL support, math functions, and time measurement tools.
-



2. Main Function Begins

```
constexpr size_t num_points = 10'000'000;
int inside_circle = 0;
```

- `num_points`: Total number of random points generated (10 million).
 - `inside_circle`: Counter for how many points fall inside the unit circle.
-

3. Start Timing

```
auto start = high_resolution_clock::now();
```

- Start the stopwatch to measure kernel execution time.
-

4. SYCL Queue and Device Info

```
queue q{ default_selector_v };
std::cout << "Running on: " << q.get_device().get_info<info::device::name>() << "\n";
```

- `queue q`: Submits commands to the best available SYCL-supported device (e.g., GPU or CPU).
 - Displays the device used.
-

5. Create a SYCL Buffer

```
buffer<int> count_buf(&inside_circle, 1);
```

- A buffer wrapping `inside_circle`. SYCL will handle its memory synchronization.
-

6. Launch Parallel Kernel

```
q.submit([&](handler& h) {
    auto count = count_buf.get_access<access::mode::atomic>(h);
```

- Submits a parallel kernel to the SYCL queue.
- `count` is an **atomic accessor**, allowing safe concurrent updates to `inside_circle`.

7. Monte Carlo Computation

```
h.parallel_for(range<1>{num_points}, [=](id<1> i) {
    size_t id = i[0];
    float x = std::fmod(std::abs(std::sin(id * 12.9898f) * 43758.5453f), 1.0f) * 2.0f - 1.0f;
    float y = std::fmod(std::abs(std::cos(id * 78.233f) * 12345.6789f), 1.0f) * 2.0f - 1.0f;

    if (x * x + y * y <= 1.0f)
        count[0].fetch_add(1);
});
```

- **Random number generation** is done using deterministic formulas (avoiding random engines for parallel efficiency).
- Each point (x, y) lies between -1 and 1.
- If it's inside the unit circle ($x^2 + y^2 \leq 1$), increment the counter.

8. Wait for Kernel Completion

```
q.wait();
```

- Ensures all submitted work is done before continuing.

9. Estimate π

```
float pi = 4.0f * inside_circle / num_points;
std::cout << "Estimated Pi = " << pi << "\n";
```

- Ratio of points inside the circle approximates $\pi/4$.
- Multiply by 4 to get the final estimate.

10. Stop Timing and Display

```
auto end = high_resolution_clock::now();
auto duration = duration_cast<milliseconds>(end - start);
std::cout << "Time taken: " << duration.count() << " ms\n";
```

- Stops the stopwatch and prints how long the whole process took.
-

✓ Summary:

This SYCL program runs a **parallel Monte Carlo simulation** to estimate π using **10 million points**. It:

- Uses a SYCL **queue** to offload work to a GPU/CPU.
- Computes all points in parallel.
- Uses **deterministic random number generation** with trigonometric functions.
- Uses **atomic operations** to avoid race conditions.

Want me to modify this code to use real random numbers or show the same simulation in CUDA/OpenMP?